

# **Machine Learning**

## **Lecture 9: Convolutional Neural Networks**

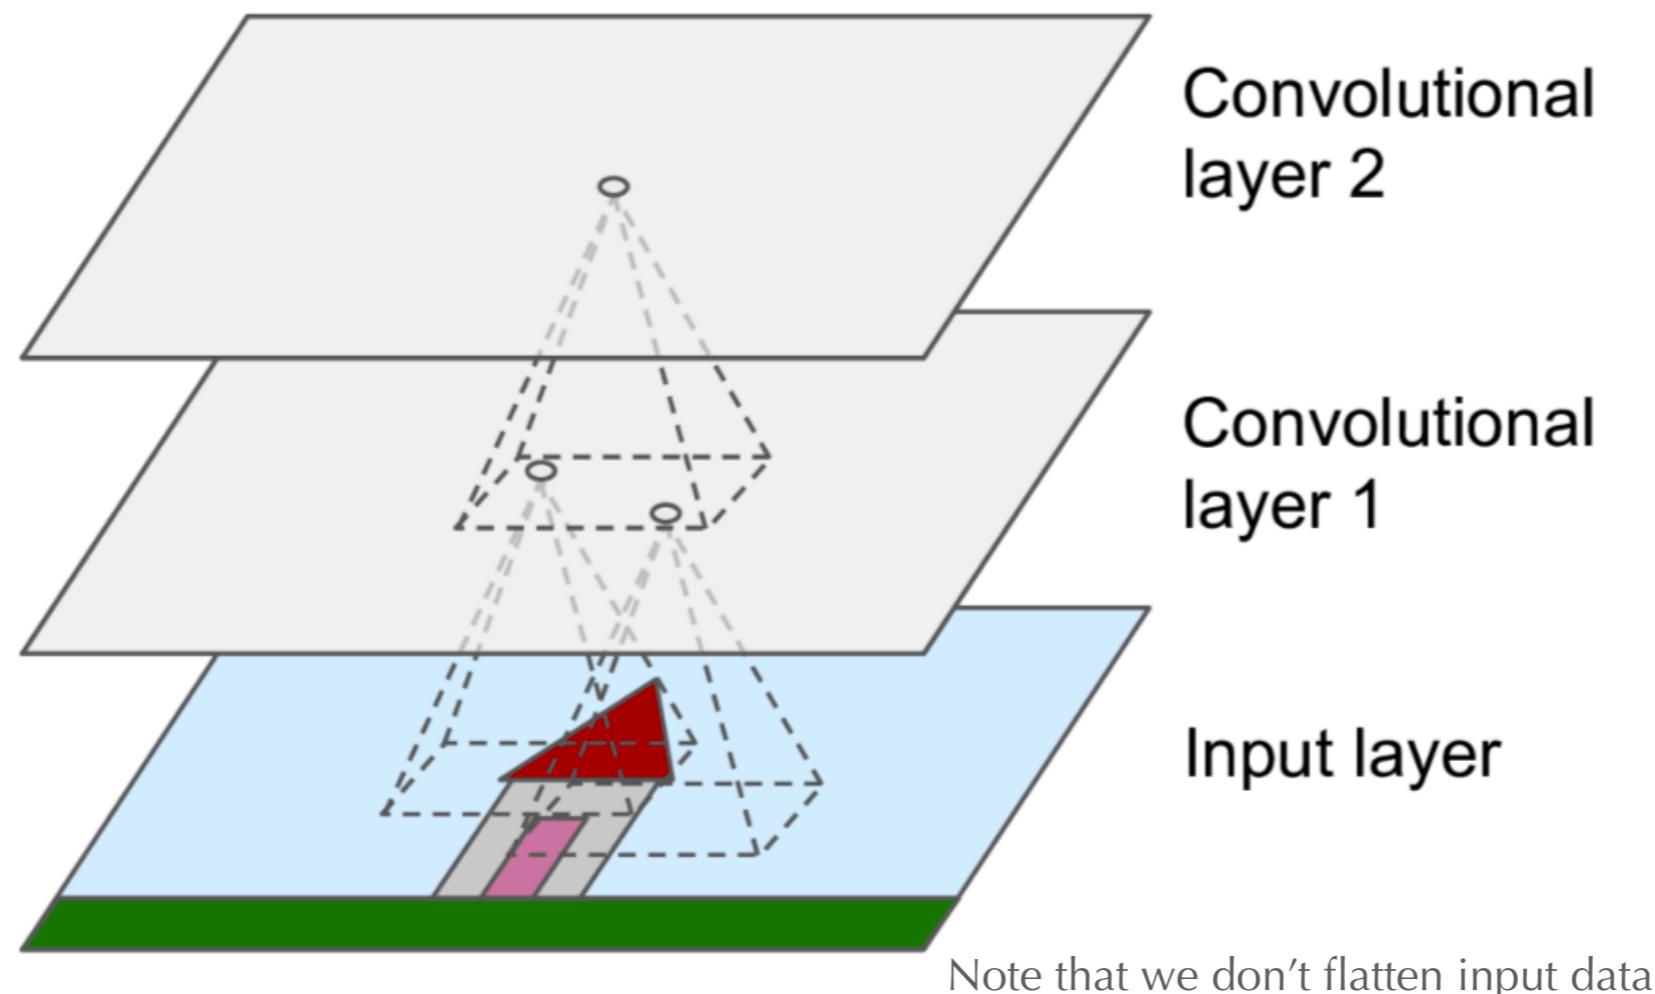
**Instructor: Dr. Farhad Pourkamali Anaraki**

# Overview

- Thanks to the increase in computational power and the amount of available training data, convolutional neural networks (CNNs) have achieved great performance on complex visual tasks
  - Image search services, self-driving cars, video classification systems, etc.
  - Not restricted to visual applications, e.g., voice recognition
- We explain building blocks of CNNs and how to implement them using Keras
  - Convolutional layers
  - Pooling layers
  - CNN architectures

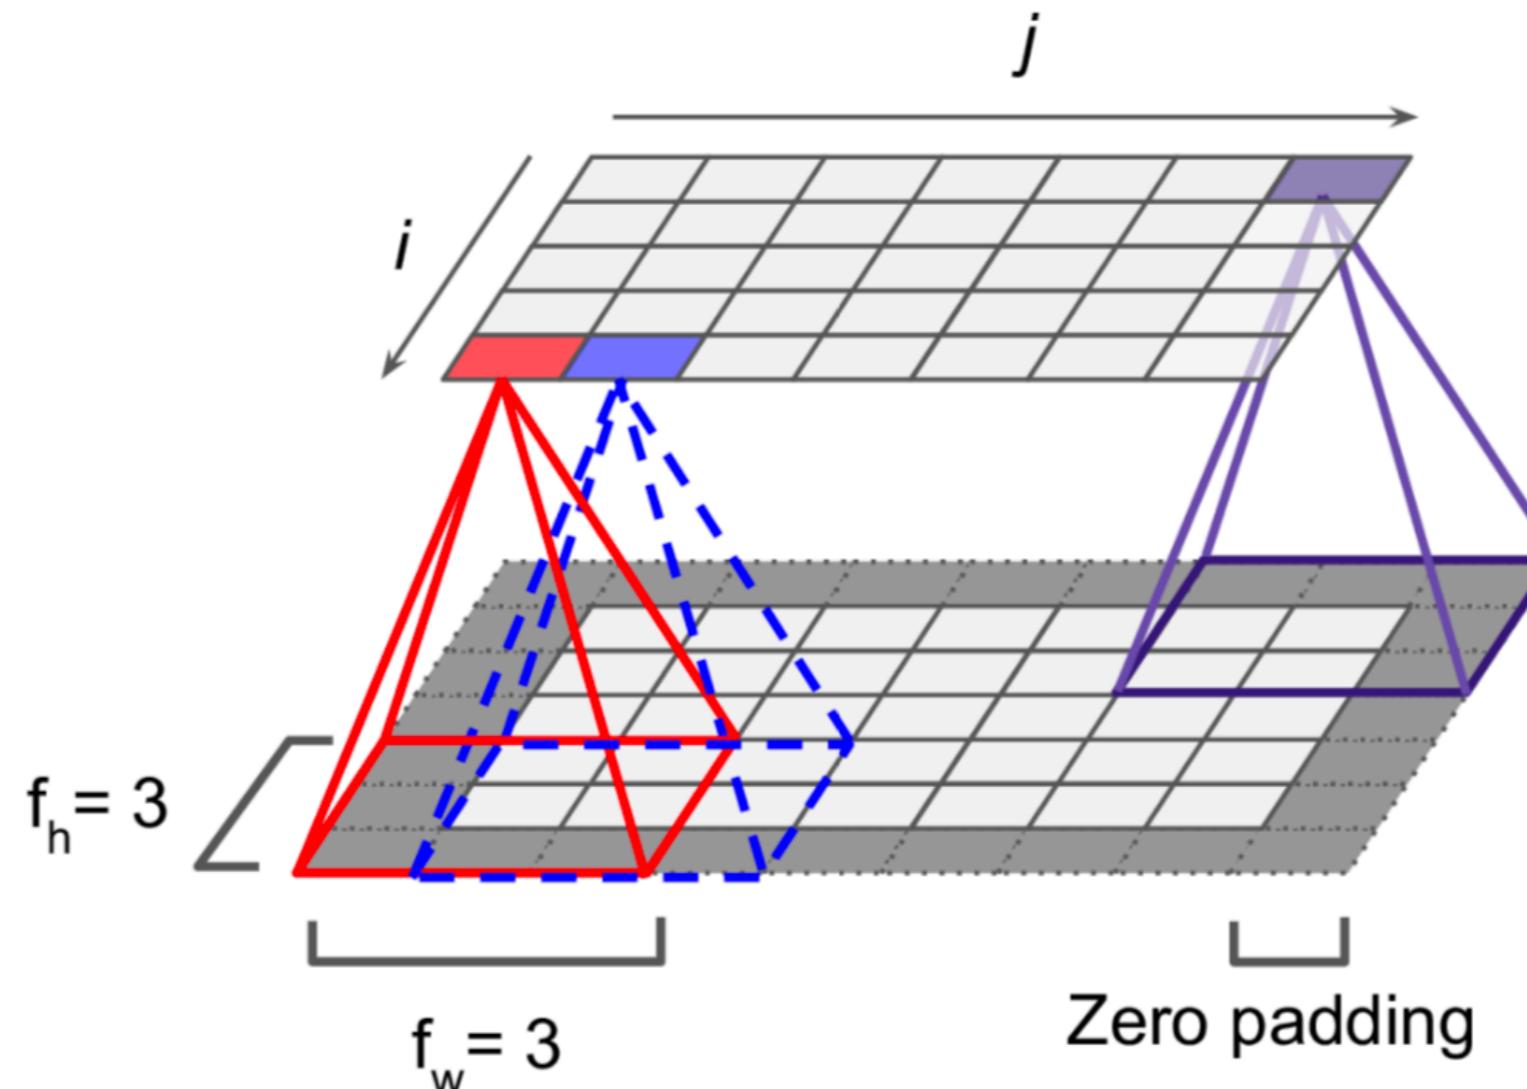
# Convolutional layers

- Neurons in the first *convolutional layer* are **not** connected to every single pixel in the input image, but only to pixels in their receptive fields
- Each neuron in the second *convolutional layer* is connected only to neurons located within a small rectangle in the first layer
- Hierarchical structure: low-level features to high-level features



# Convolutional layers

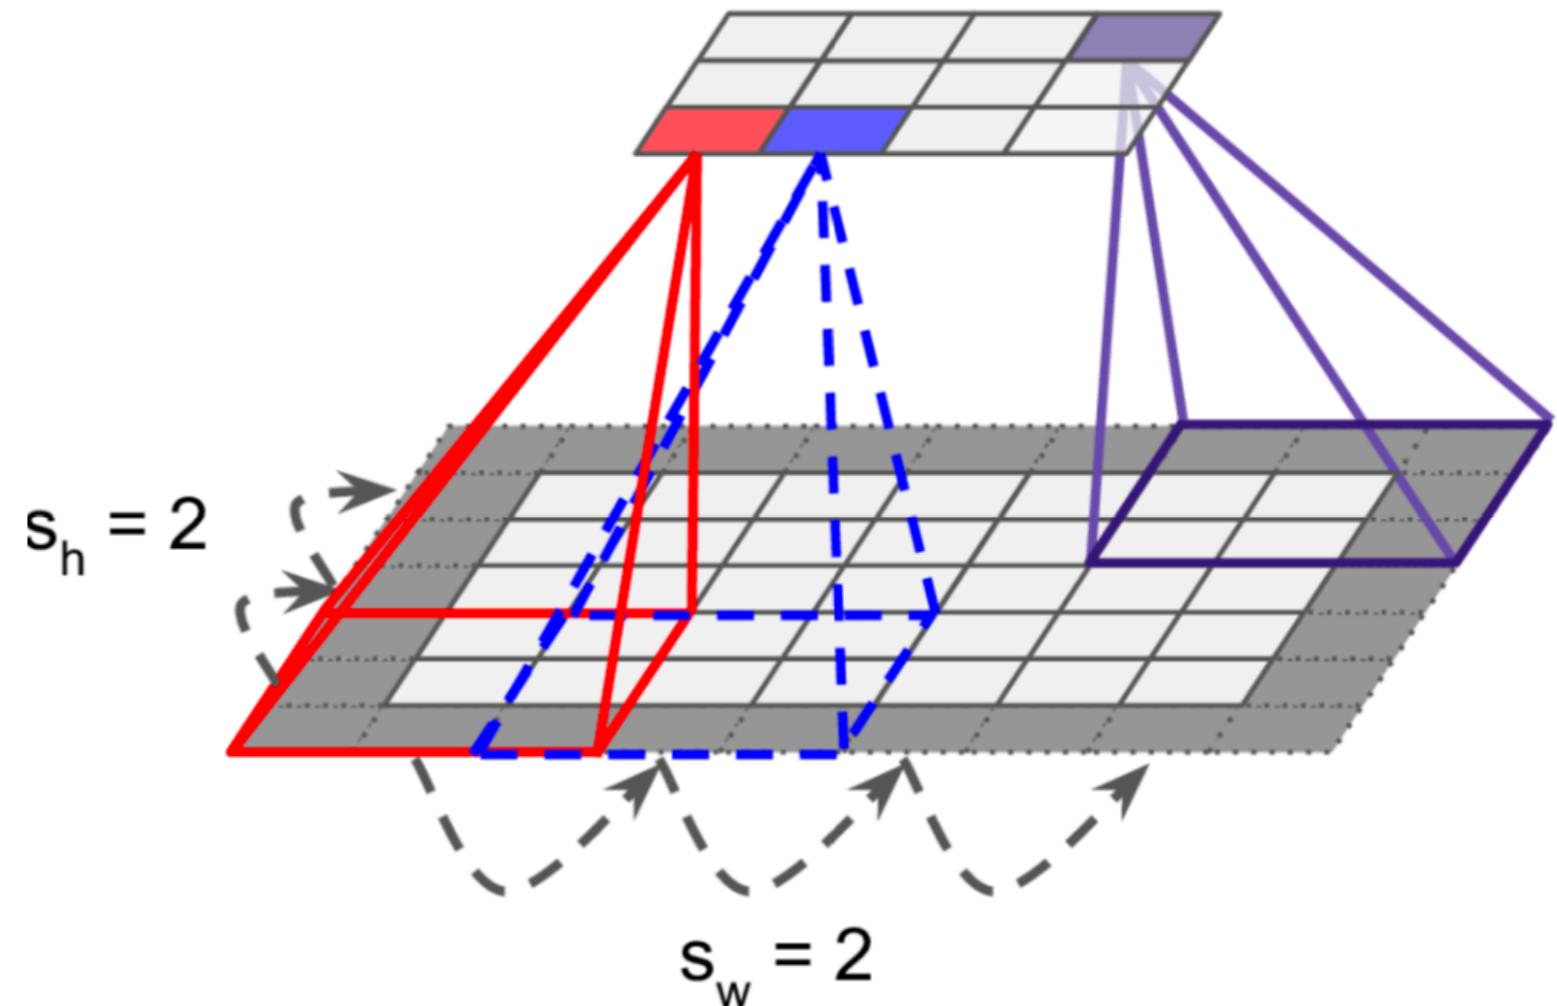
- To be more formal, each neuron in row  $i$  and column  $j$  of a given layer is connected to the outputs of the neurons in the previous layer located in rows  $i$  to  $i + f_h - 1$  and columns  $j$  to  $j + f_w - 1$ , where  $f_h$  and  $f_w$  are the height and width of the receptive field



- We add zeros around the inputs to have the same number of rows and columns

# Stride

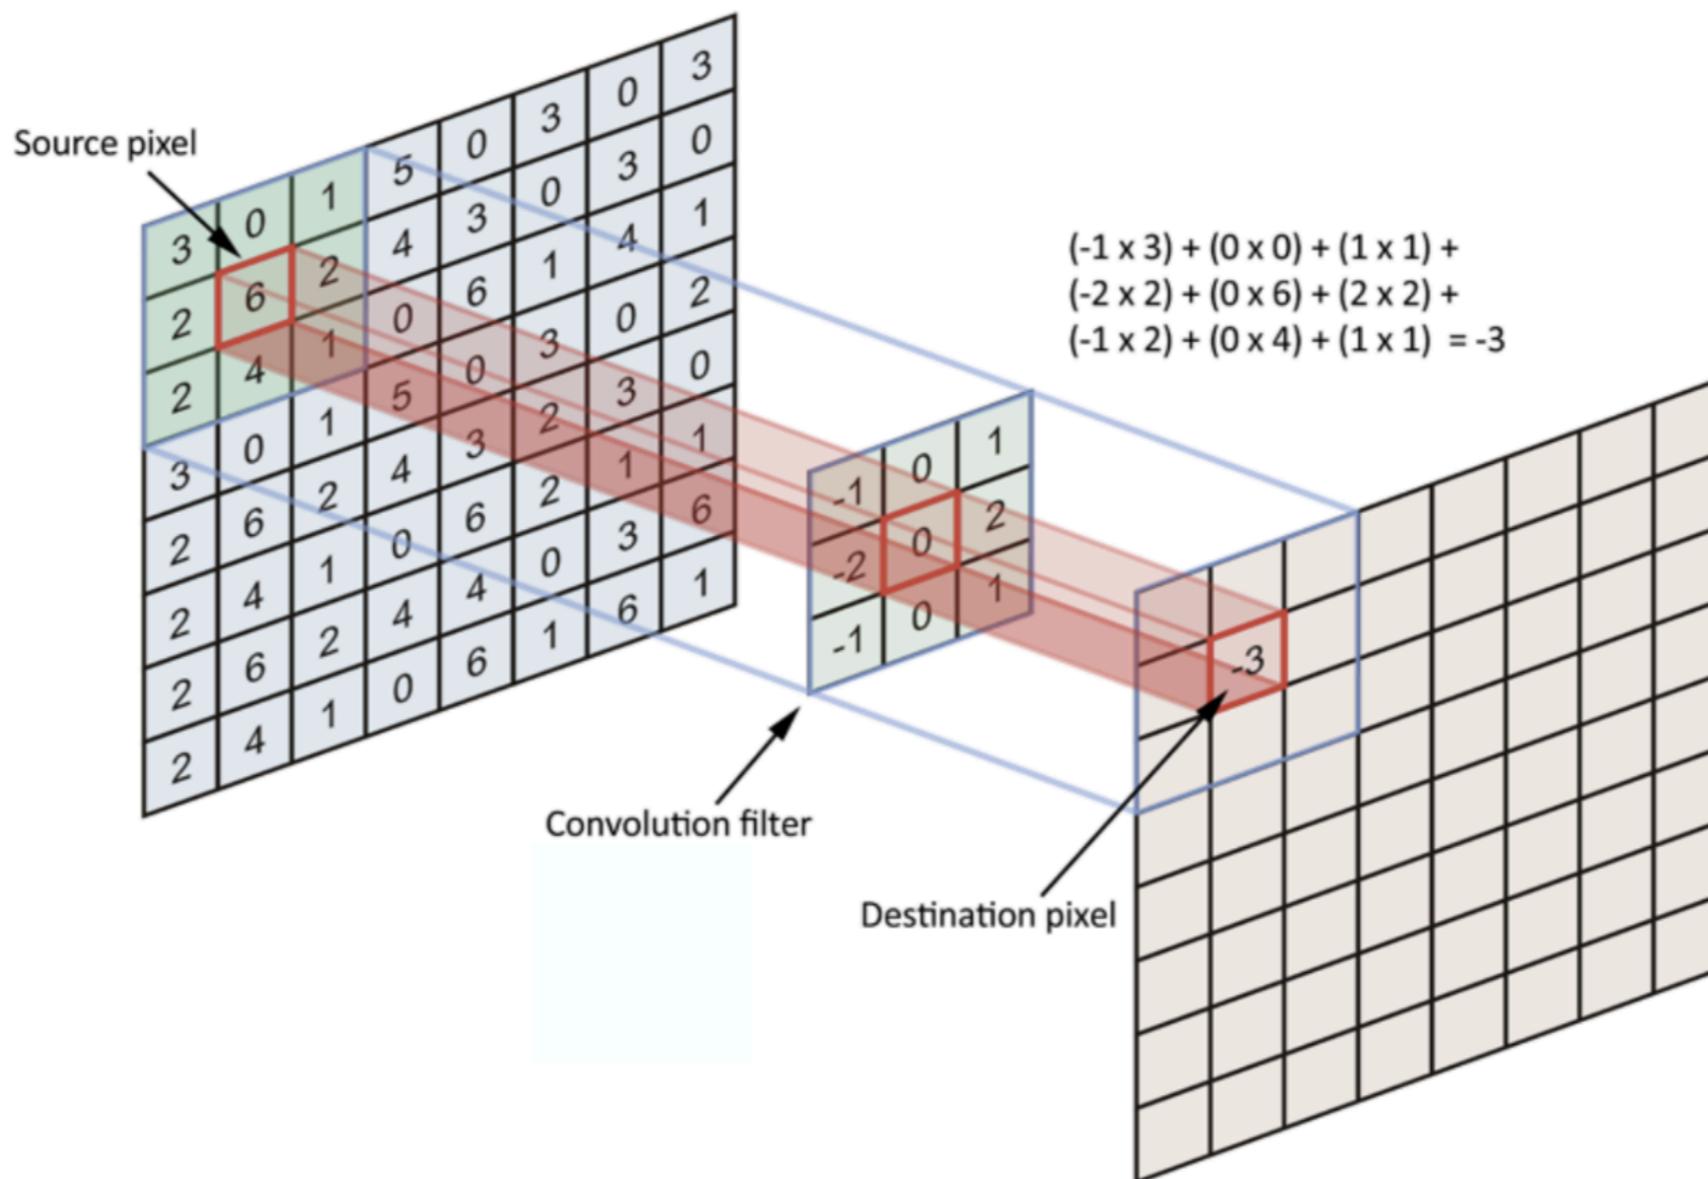
- The shift from one receptive field to the next is called the **stride**
- 3x3 receptive fields and a stride of 2



- Neuron located in row  $i$  and column  $j$  is connected to rows  $i \times s_h$  to  $i \times s_h + f_h - 1$  and columns  $j \times s_w$  to  $j \times s_w + f_w - 1$

# Filters

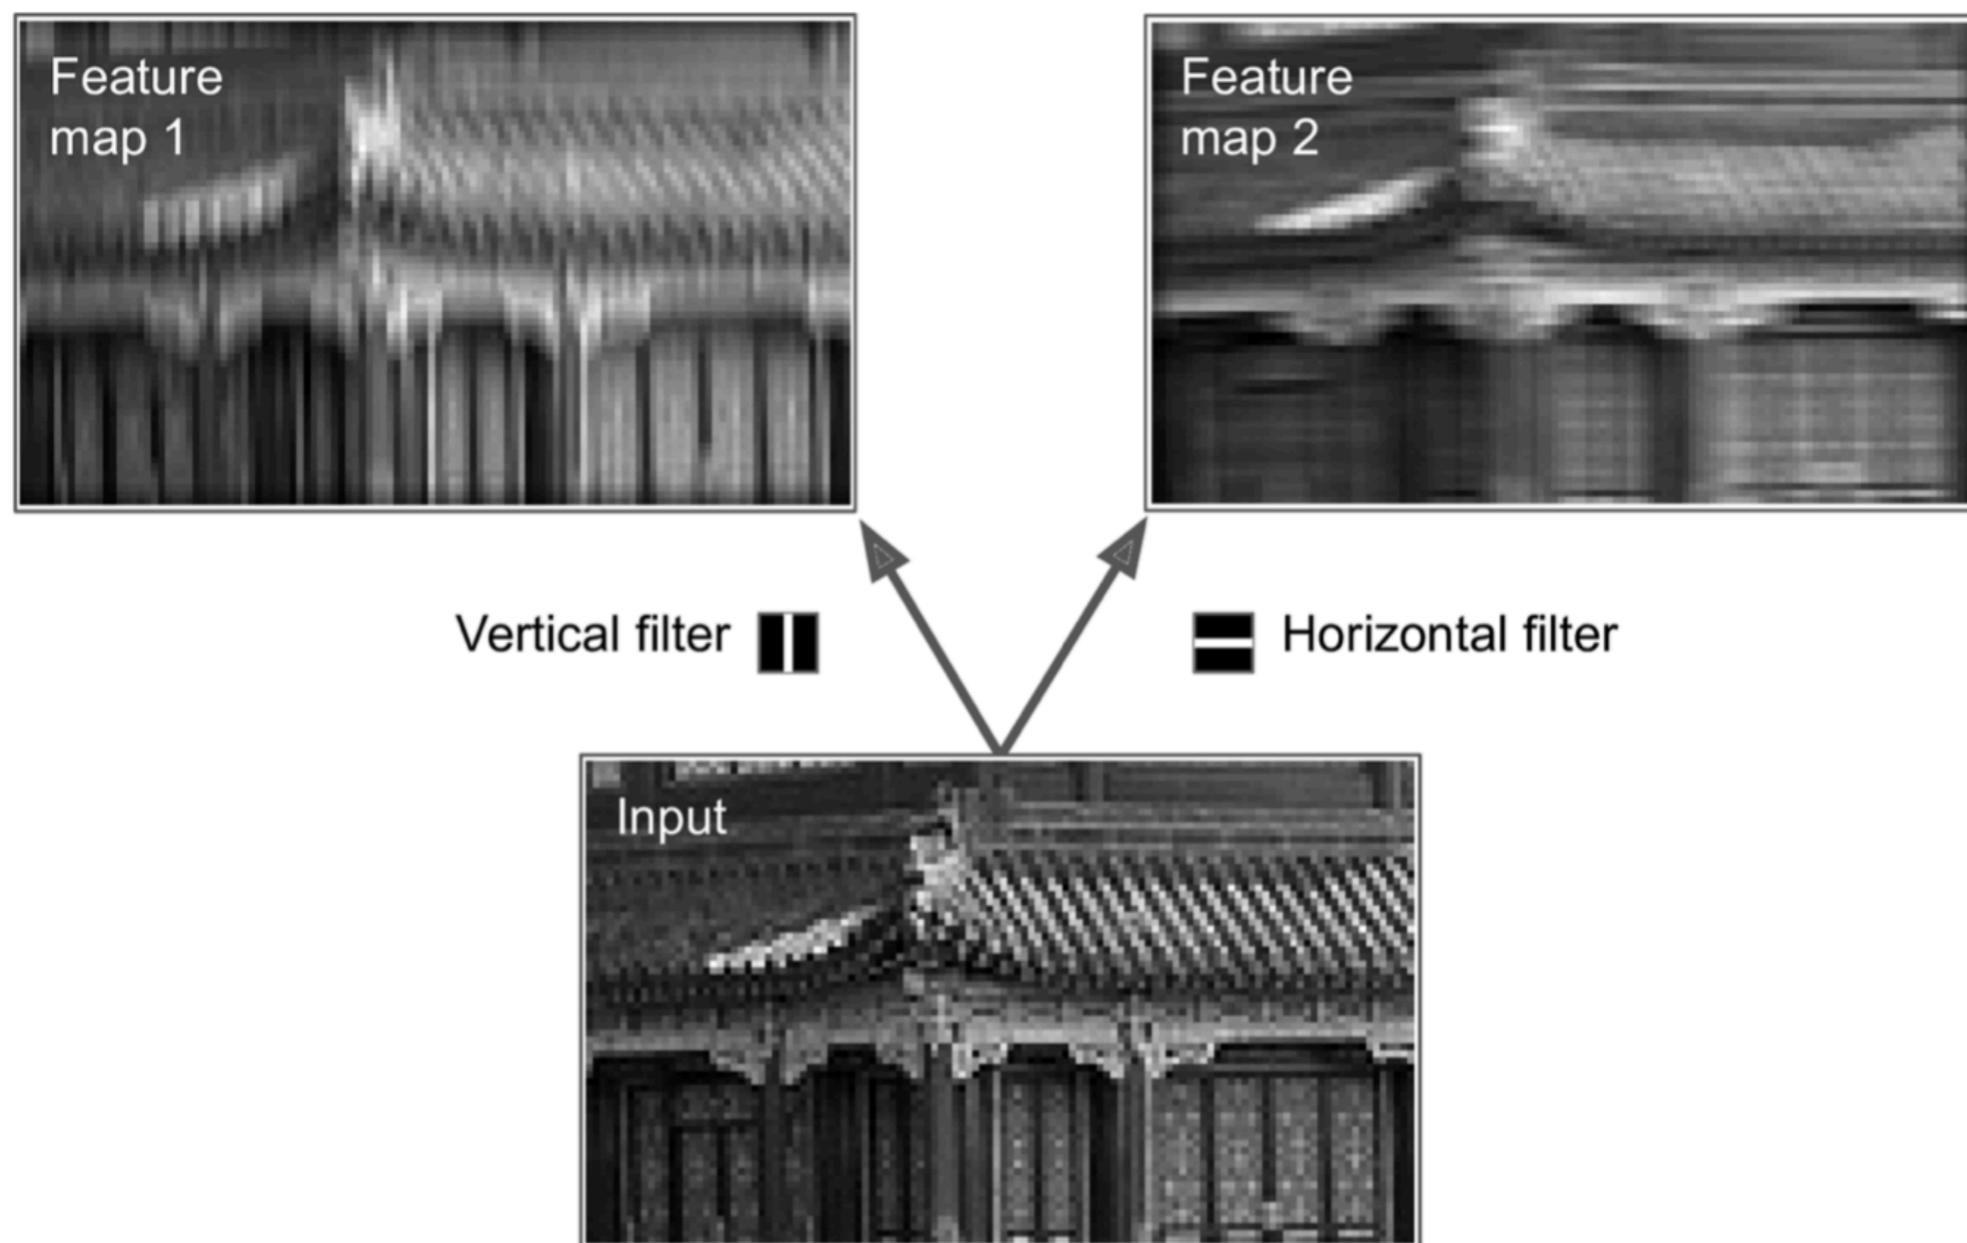
- A neuron's weight can be represented as a small image of the size of the receptive field
- We refer to sets of weights as filters or convolutional kernels



The convolution operation.

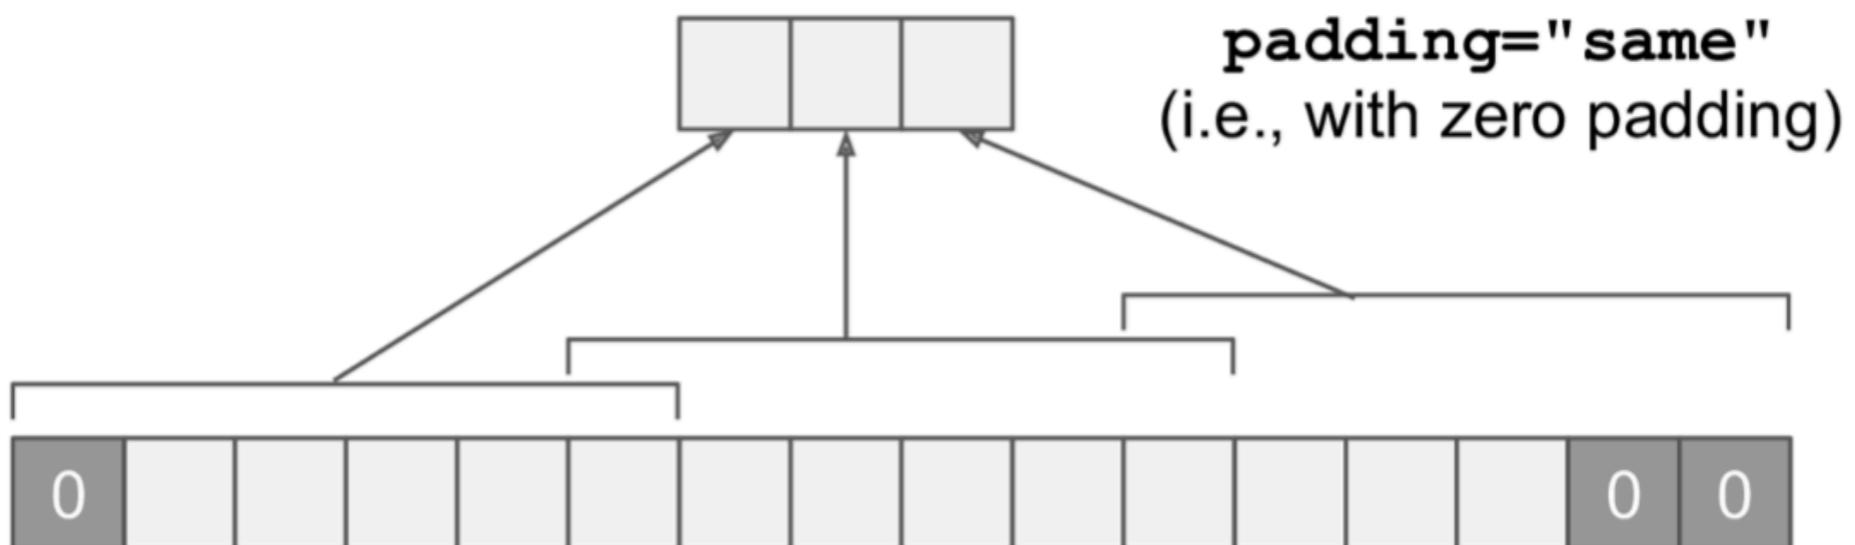
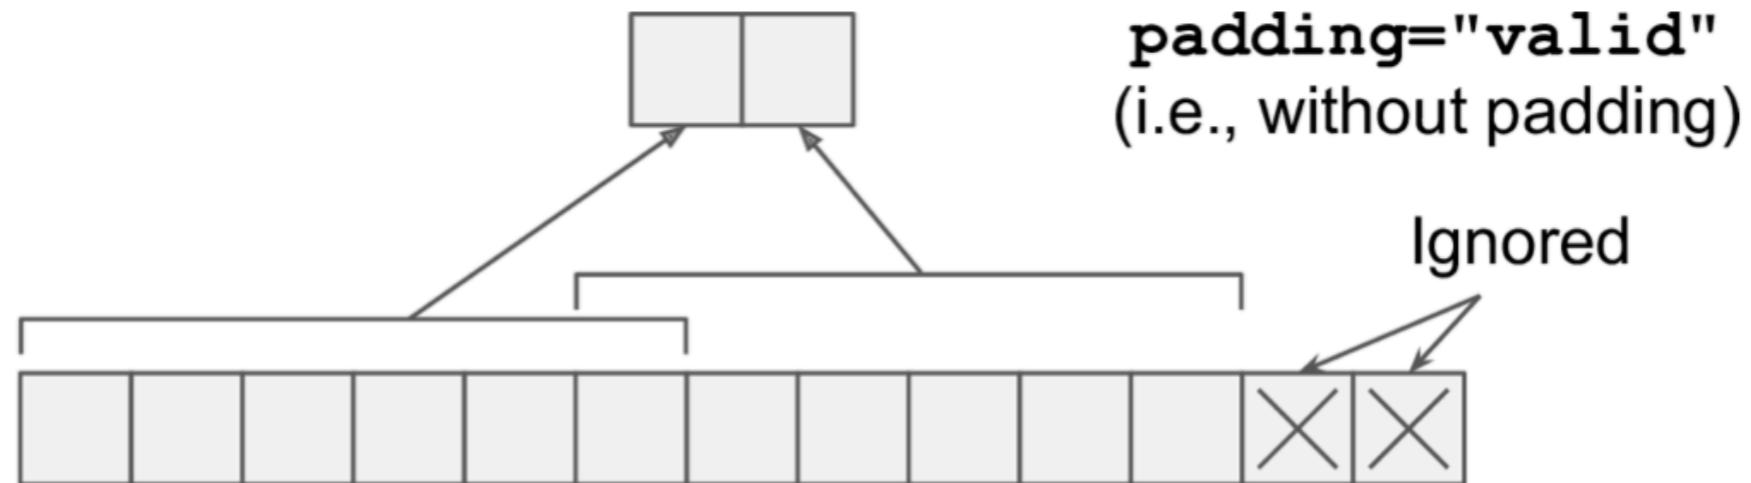
# Examples

- Let us consider two possible sets of weights, i.e., filters
- We obtain feature maps, highlighting the areas in an image that activate the filter the most



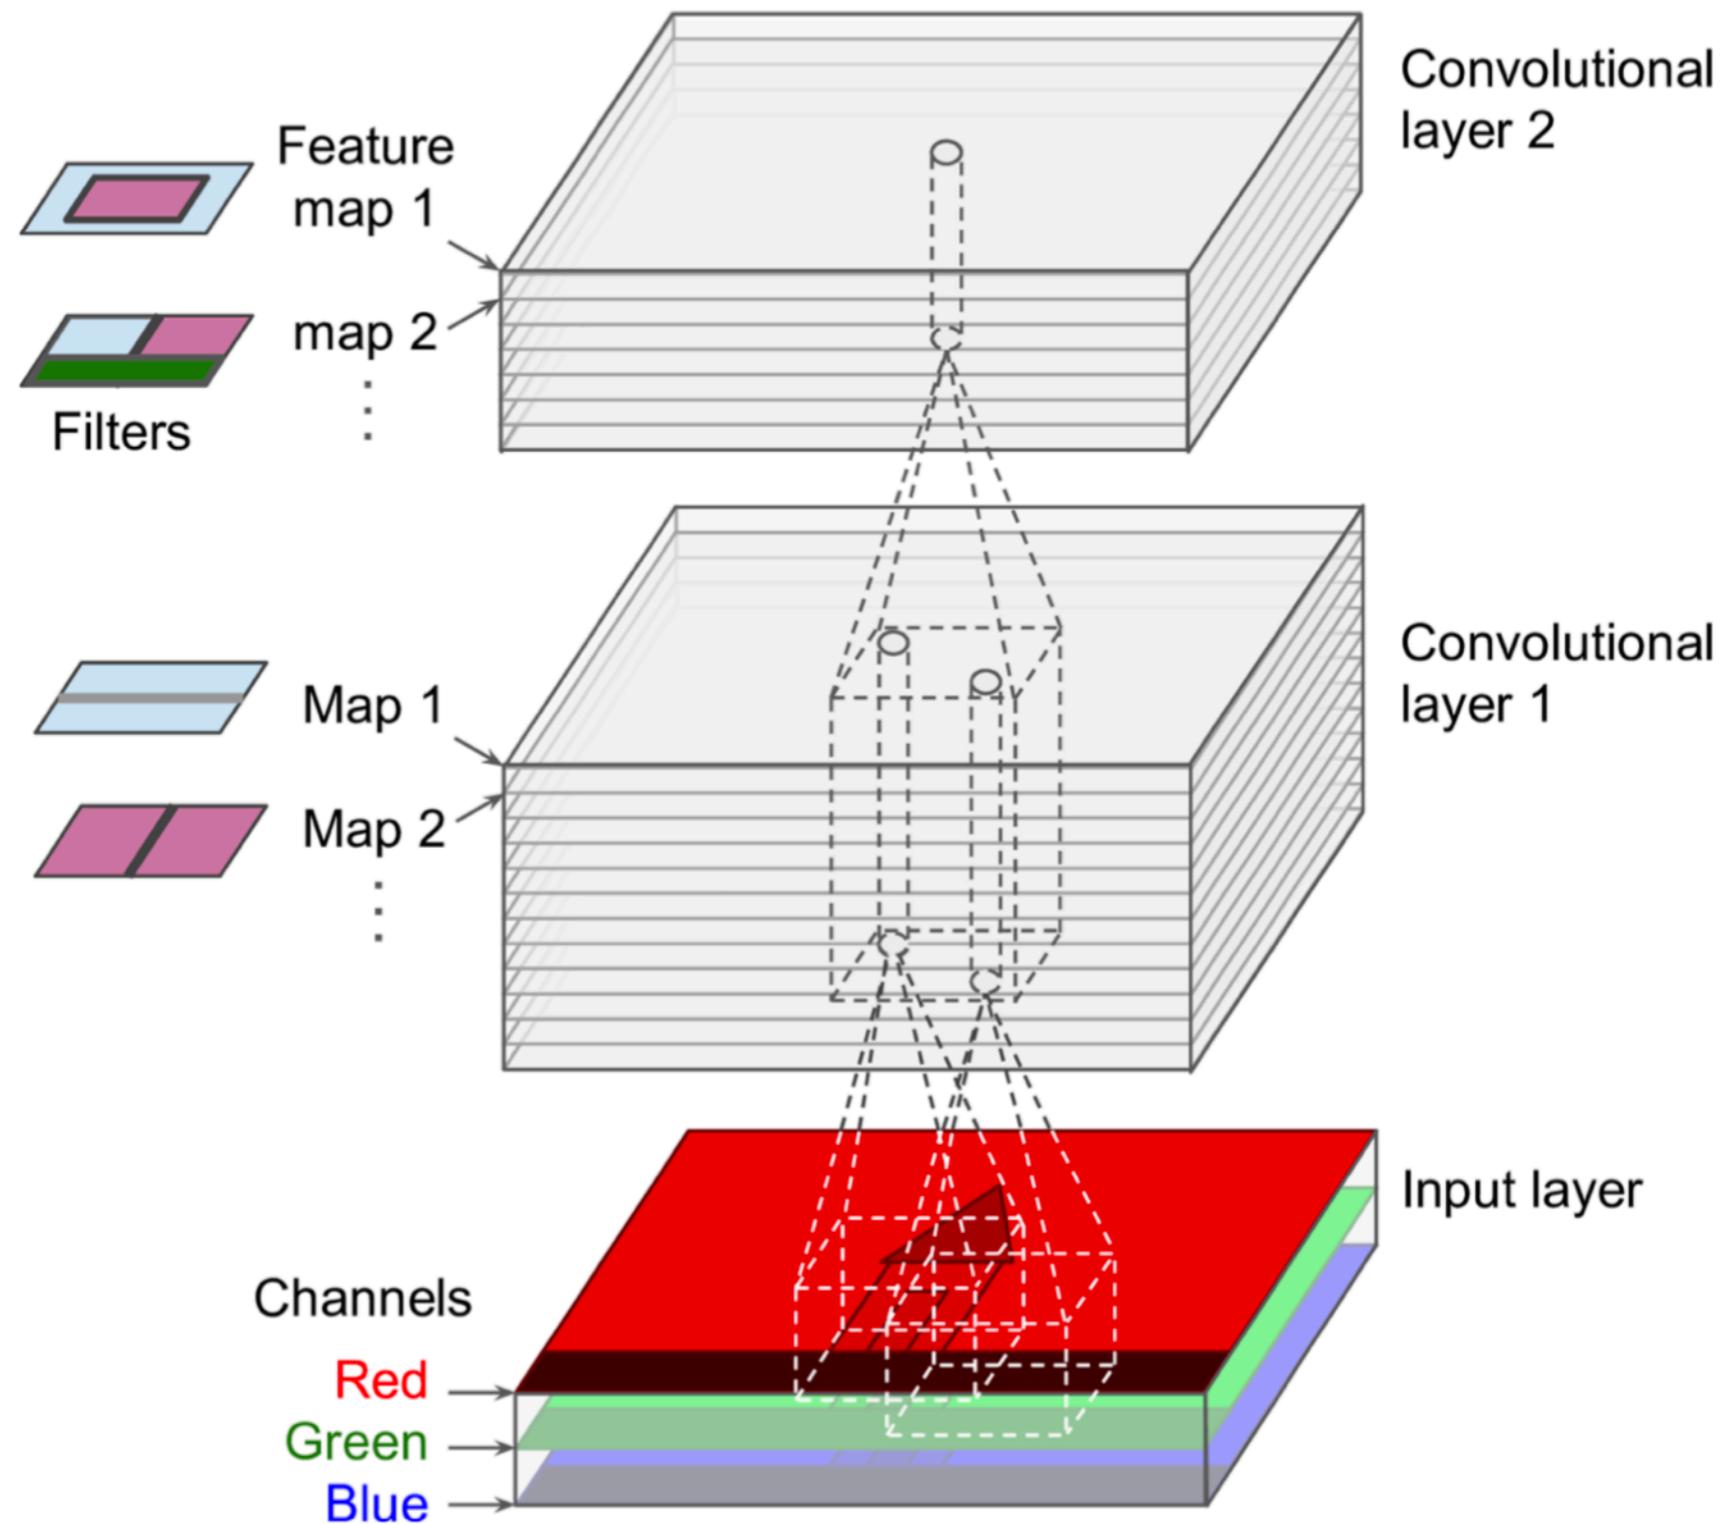
# Two different kinds of padding

- “SAME” vs “VALID” padding (input width=13, filter width=6, stride=5)



# Stacking multiple feature maps

- The number of parameters in CNNs is dramatically reduced because all neurons in a feature map share the same parameters for each filter



# Mathematical representation

- We calculate the weighted sum of all the inputs, plus the bias term

*Equation 14-1. Computing the output of a neuron in a convolutional layer*

$$z_{i,j,k} = b_k + \sum_{u=0}^{f_h-1} \sum_{v=0}^{f_w-1} \sum_{k'=0}^{f_{n'}-1} x_{i',j',k'} \times w_{u,v,k',k} \quad \text{with} \quad \begin{cases} i' = i \times s_h + u \\ j' = j \times s_w + v \end{cases}$$

- $x_{i',j',k'}$  is the output of the neuron located in layer  $l - 1$ , row  $i'$ , column  $j'$ , feature map  $k'$  (or channel  $k'$  if the previous layer is the input layer).
- $b_k$  is the bias term for feature map  $k$  (in layer  $l$ ). You can think of it as a knob that tweaks the overall brightness of the feature map  $k$ .
- $w_{u,v,k',k}$  is the connection weight between any neuron in feature map  $k$  of the layer  $l$  and its input located at row  $u$ , column  $v$  (relative to the neuron's receptive field), and feature map  $k'$ .

# Keras Implementation

- We can use the keras.layers.Conv2D

```
conv = keras.layers.Conv2D(filters=32, kernel_size=3, strides=1,  
                           padding="same", activation="relu")
```

- 32 filters each 3x3 using a stride of 1 (both horizontally and vertically) and “same” padding, and then we apply the ReLU activation function
- Number of filters, their height and width
- Strides
- Padding type

# Pooling layer

- Like convolutional layer, each neuron in a pooling layer is connected to a small rectangular receptive field. However, its goal is to aggregate the inputs using a function such as the max or mean
- The idea is to shrink the input image to reduce the computational load and memory usage

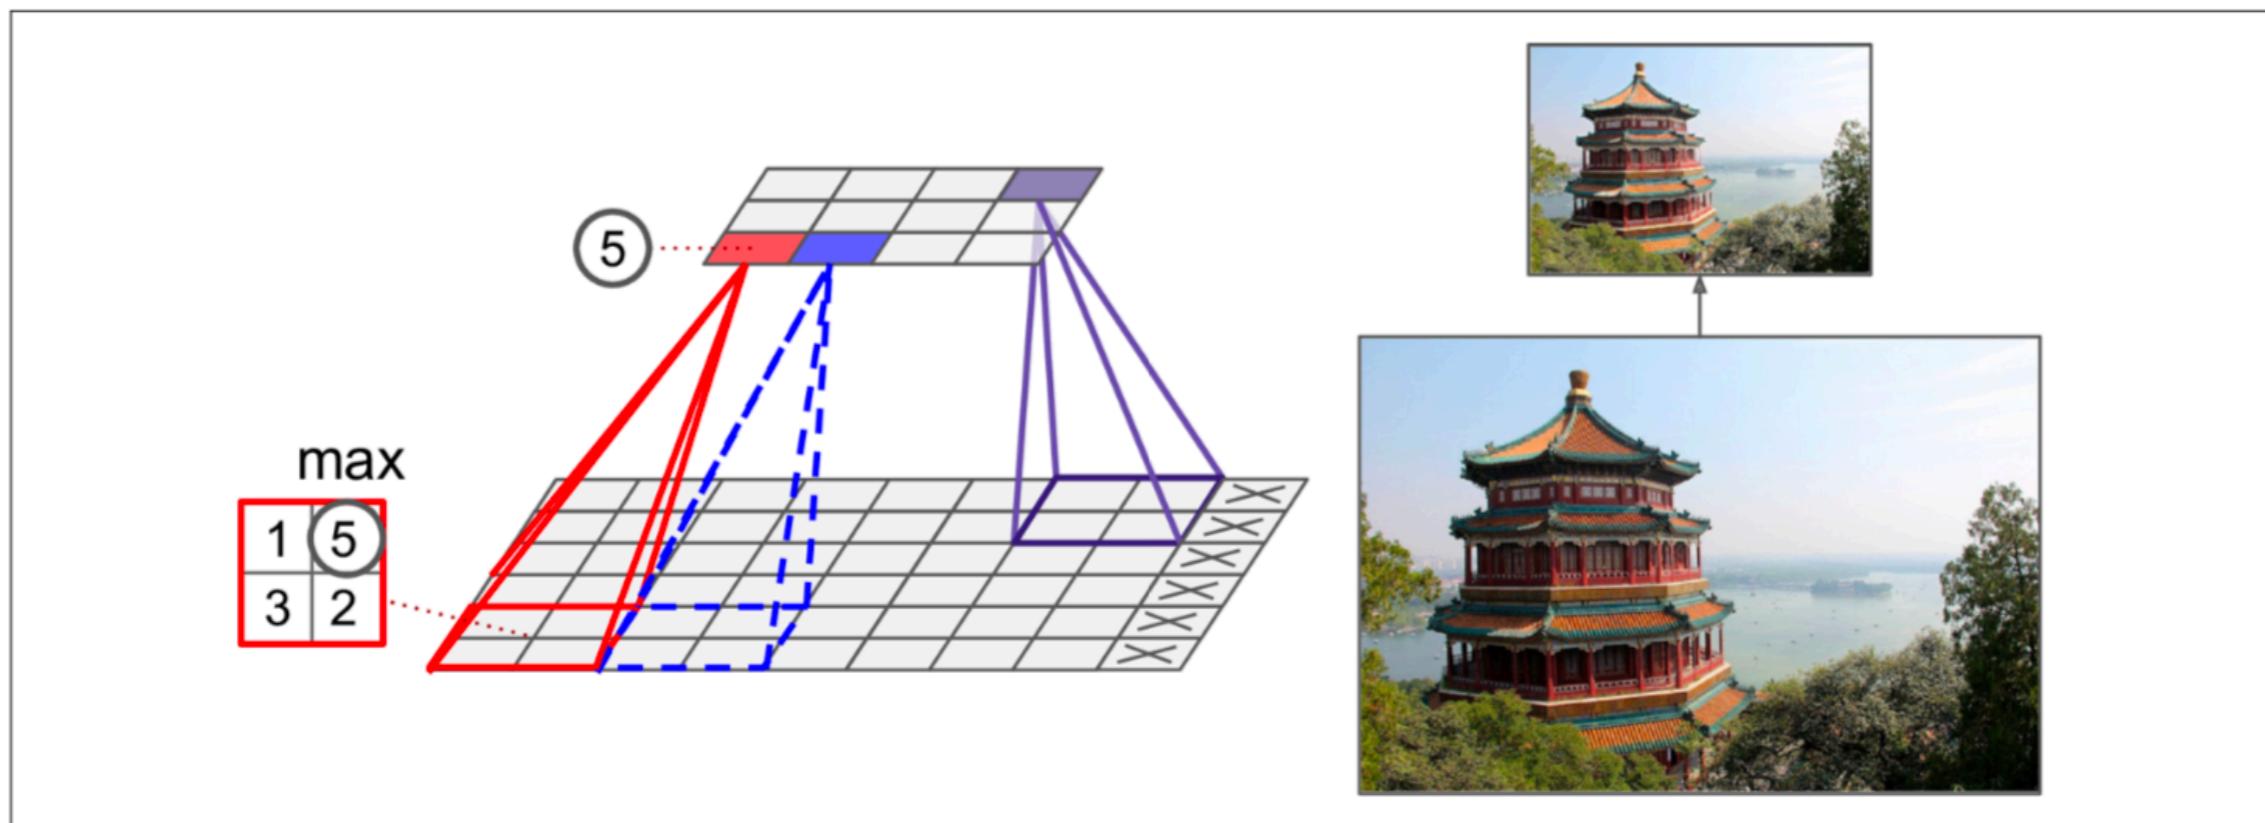
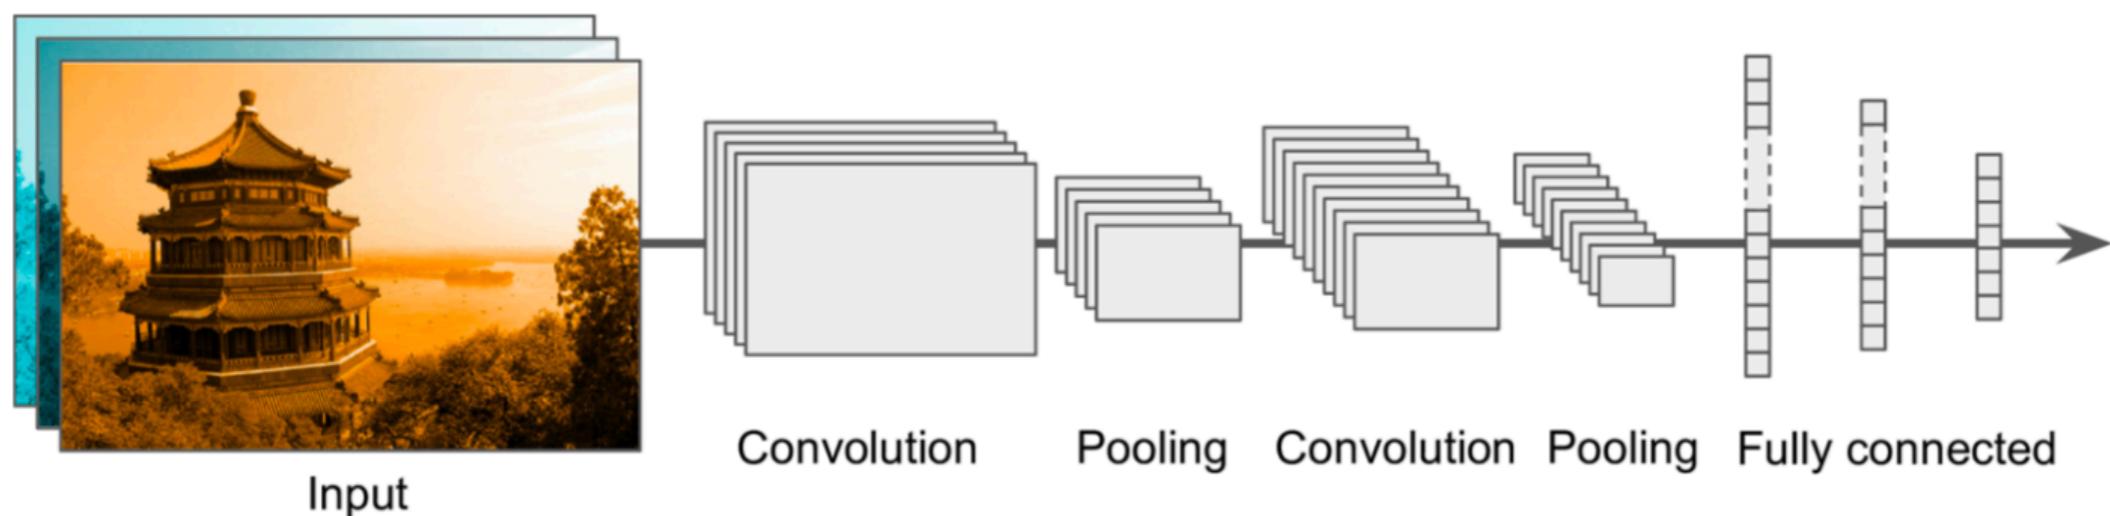


Figure 14-8. Max pooling layer ( $2 \times 2$  pooling kernel, stride 2, no padding)

```
max_pool = keras.layers.MaxPool2D(pool_size=2)
```

# CNN architectures

- Typically, a few convolutional layers, then a pooling layer, and then convolutional layers,...
- A regular feedforward neural network is added composed of fully connected layers
- Final layer outputs the prediction (e.g., a softmax layer that outputs estimated class probabilities)

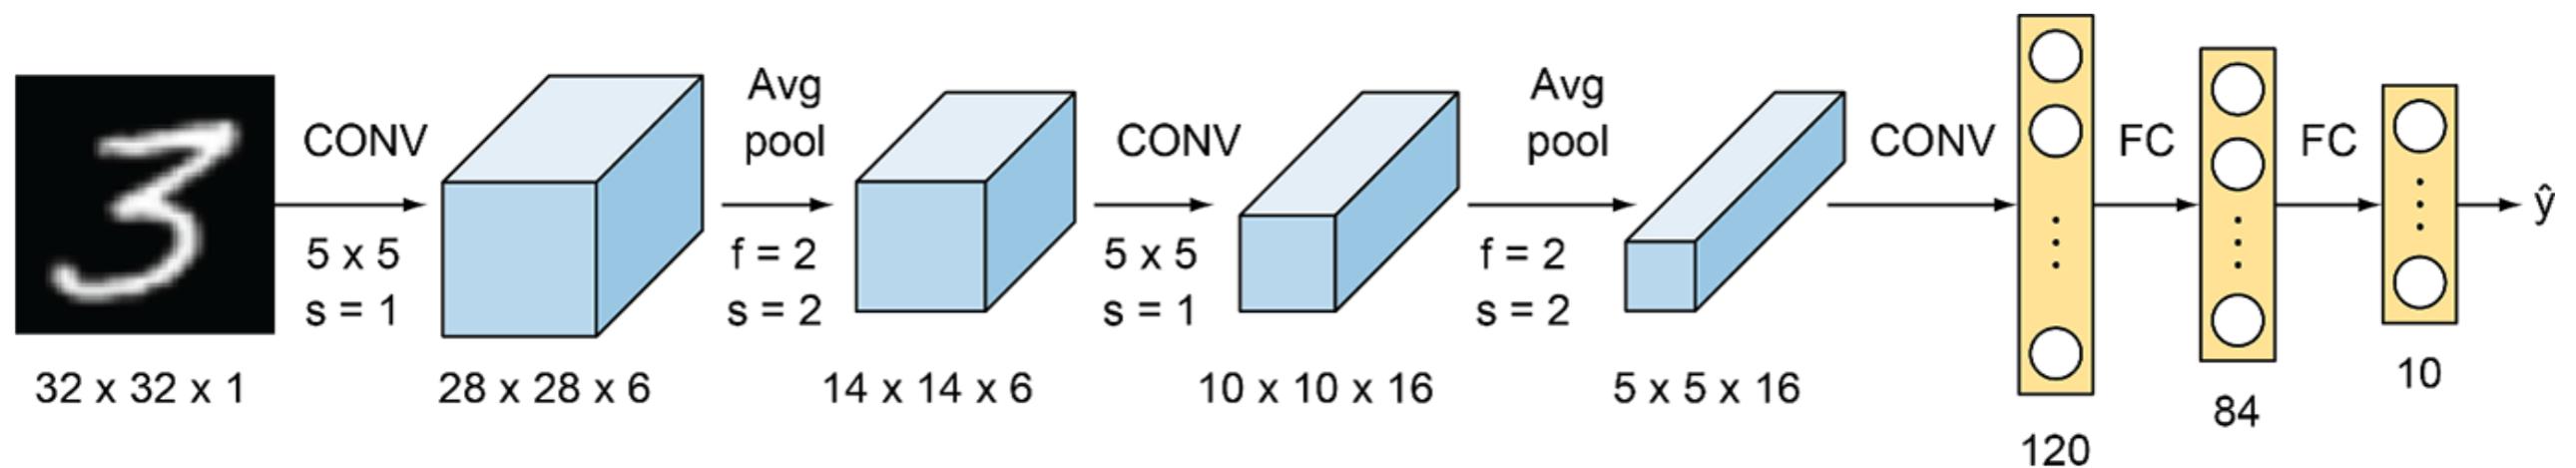


# Keras implementation

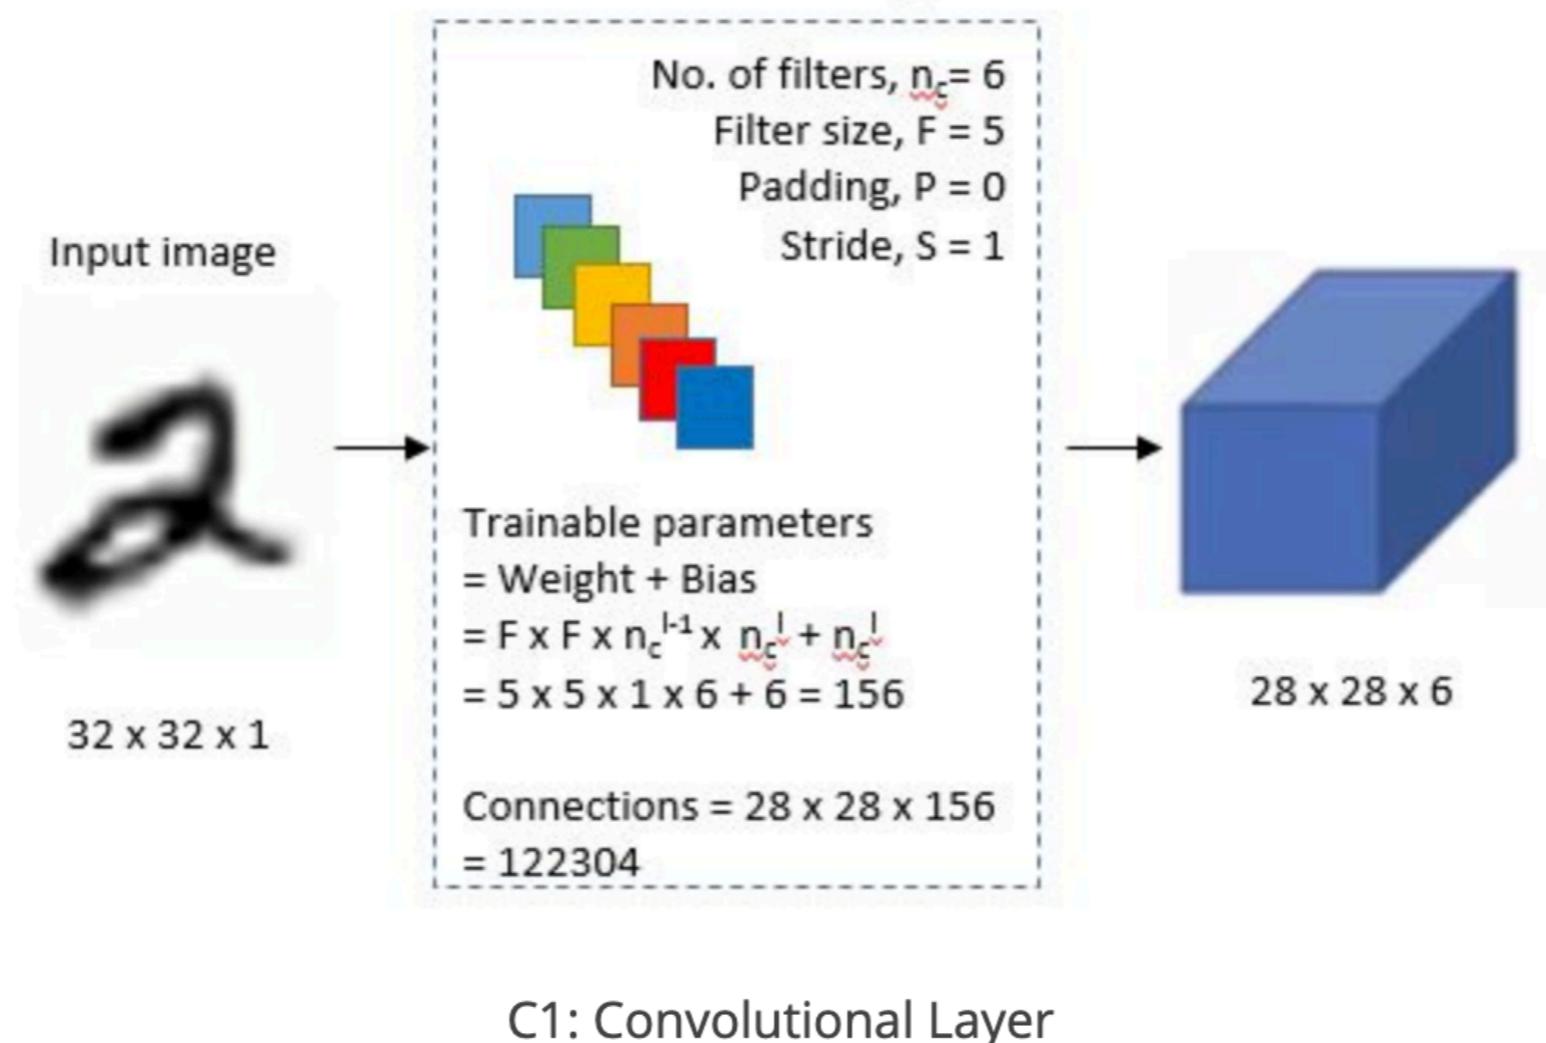
```
model = keras.models.Sequential([
    keras.layers.Conv2D(64, 7, activation="relu", padding="same",
                       input_shape=[28, 28, 1]),
    keras.layers.MaxPooling2D(2),
    keras.layers.Conv2D(128, 3, activation="relu", padding="same"),
    keras.layers.Conv2D(128, 3, activation="relu", padding="same"),
    keras.layers.MaxPooling2D(2),
    keras.layers.Conv2D(256, 3, activation="relu", padding="same"),
    keras.layers.Conv2D(256, 3, activation="relu", padding="same"),
    keras.layers.MaxPooling2D(2),
    keras.layers.Flatten(),
    keras.layers.Dense(128, activation="relu"),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(64, activation="relu"),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(10, activation="softmax")
])
```

# Popular architectures

- LeNet-5 Architecture

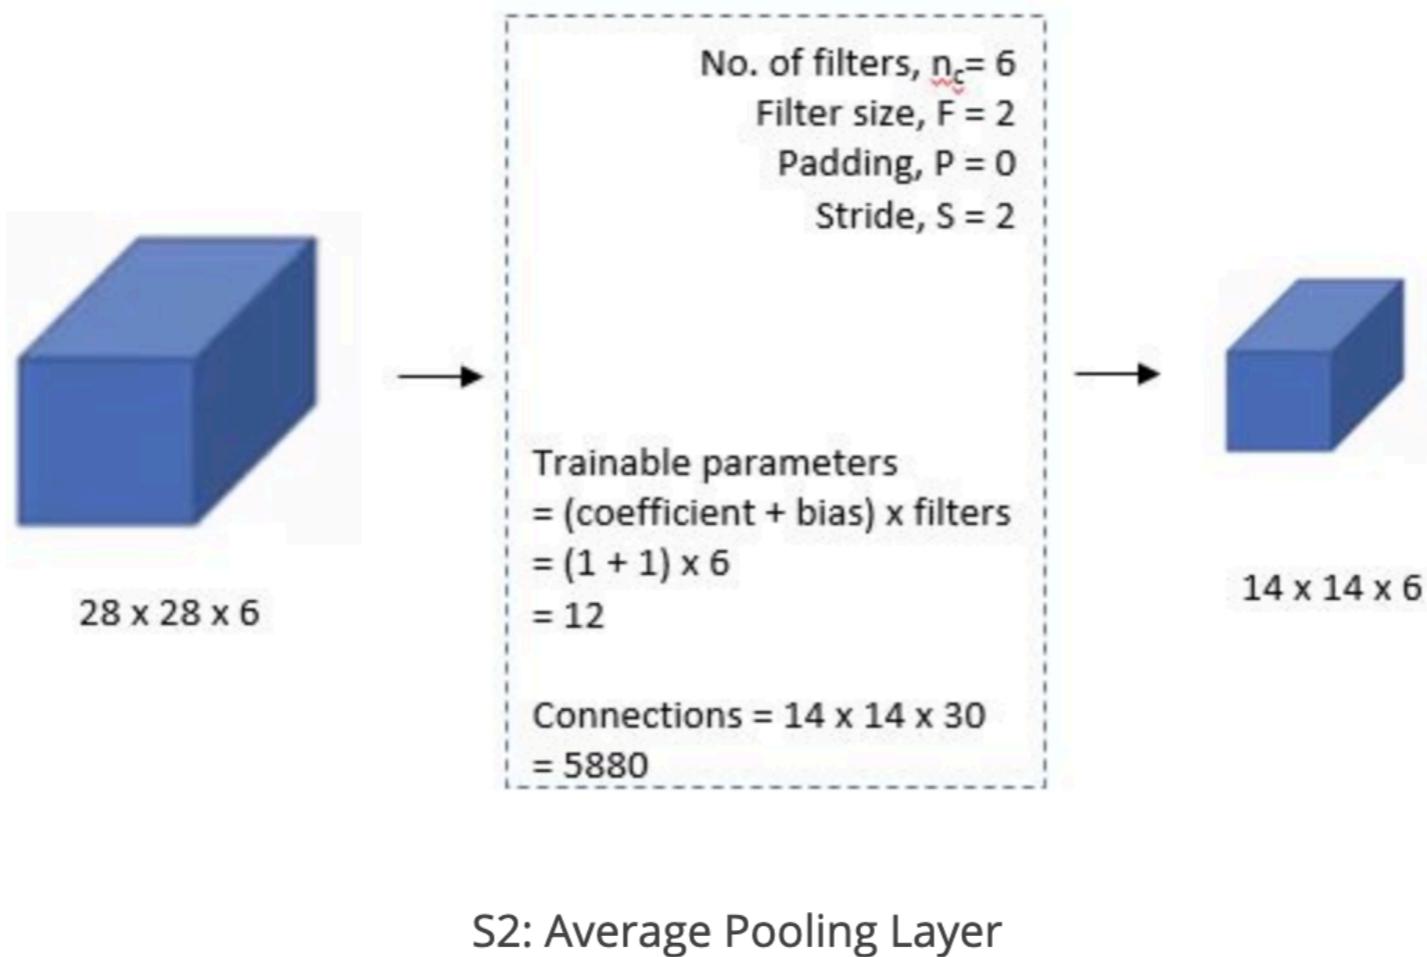


# Some known architectures

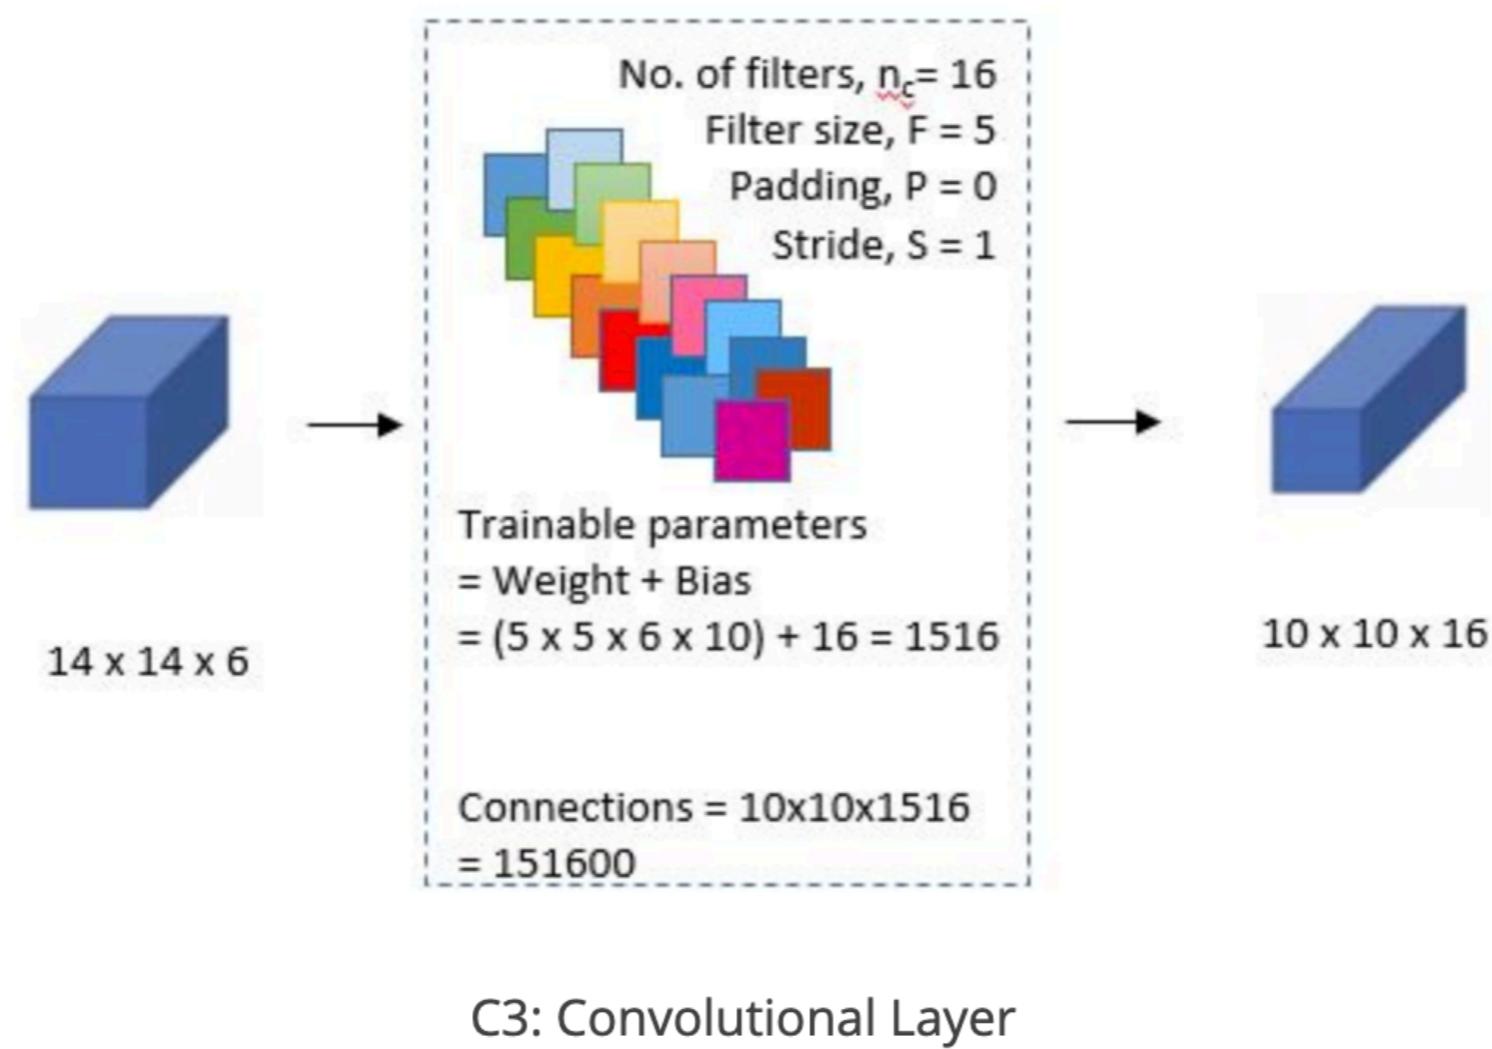


# Some known architectures

---

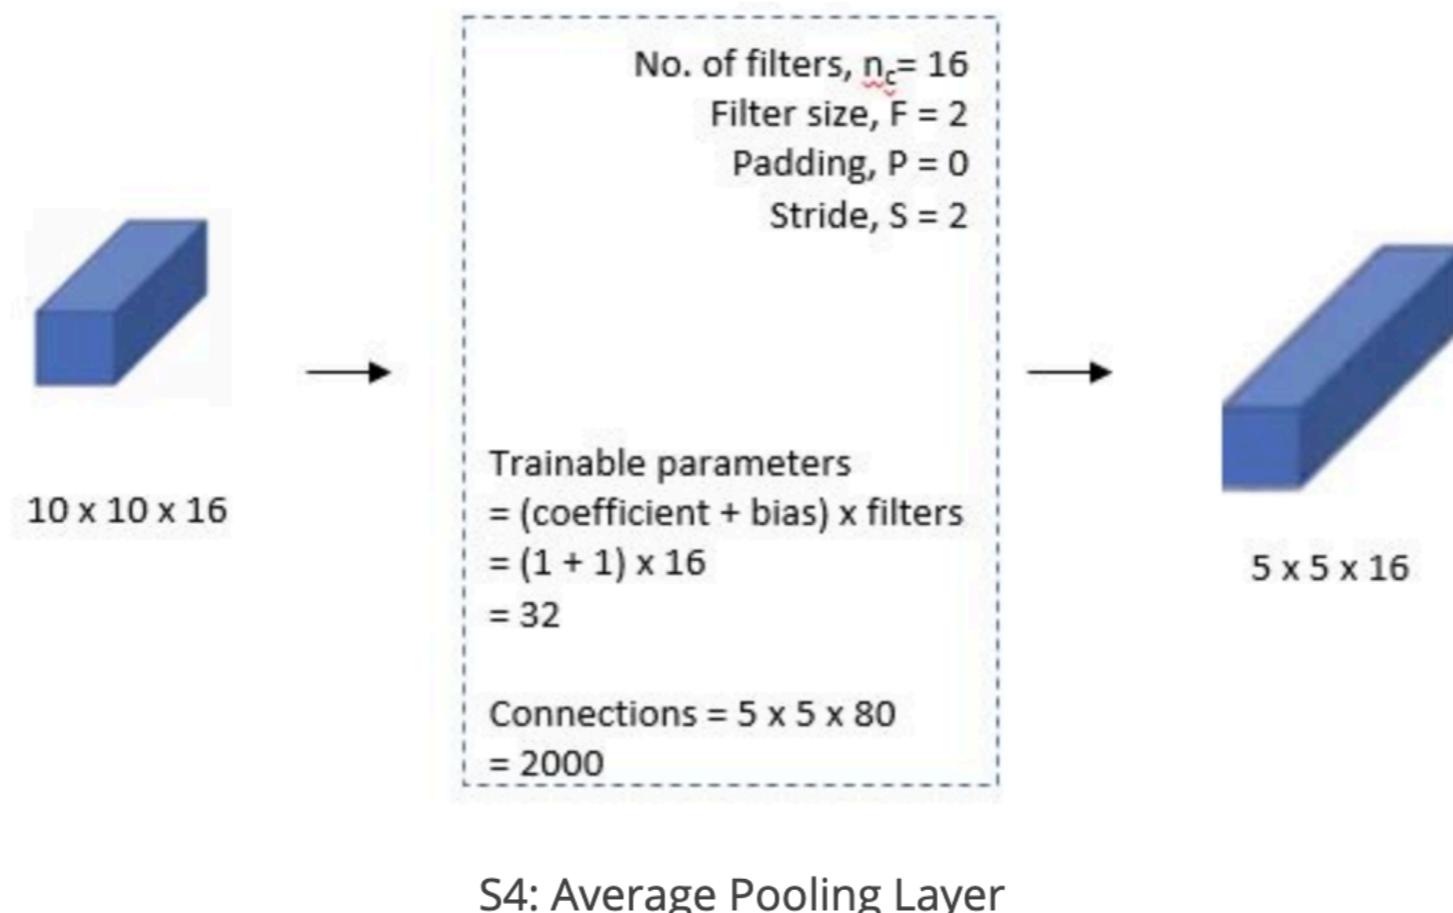


# Some known architectures



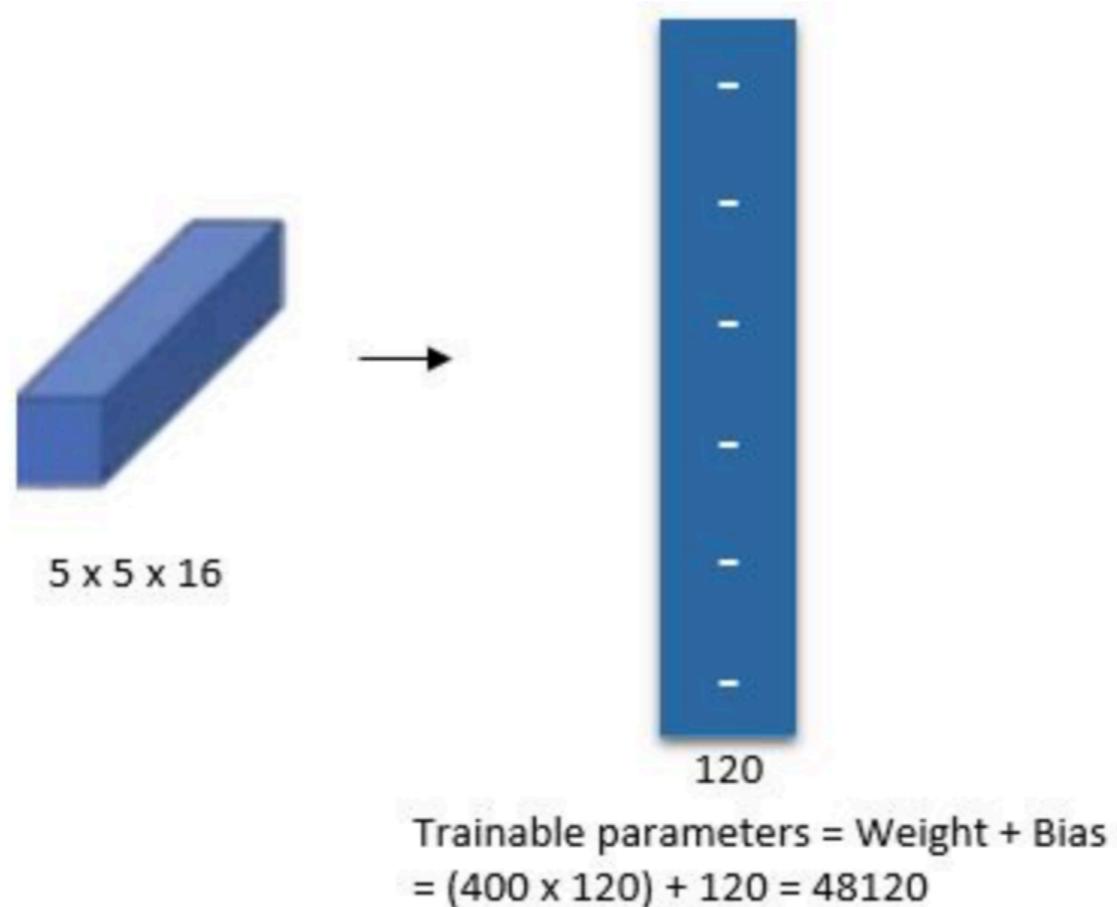
# Some known architectures

---



# Some known architectures

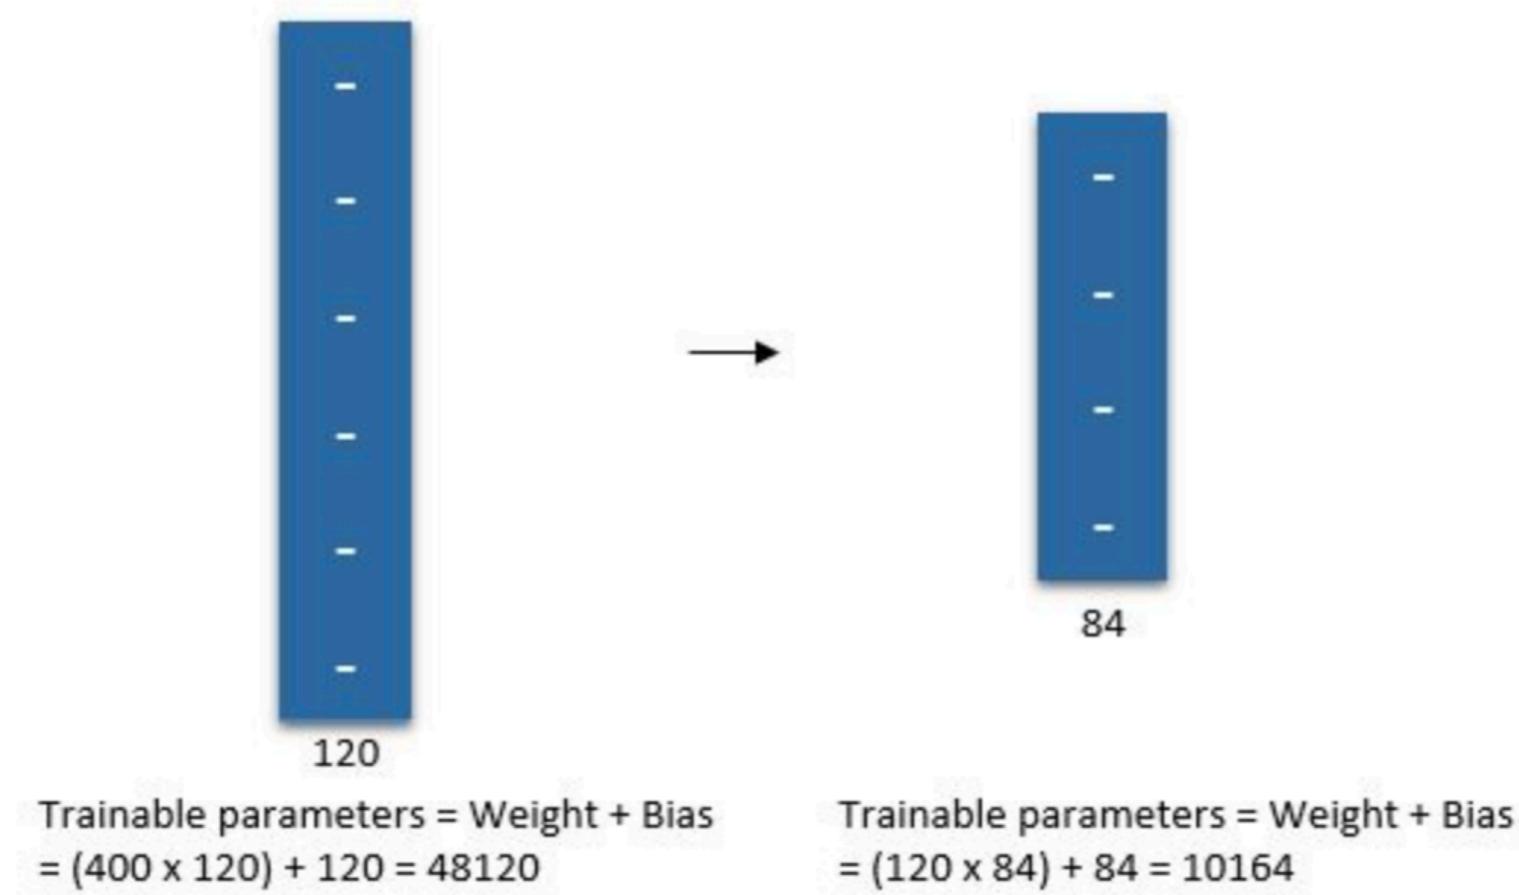
---



C5: Fully Connected Convolutional Layer

# Some known architectures

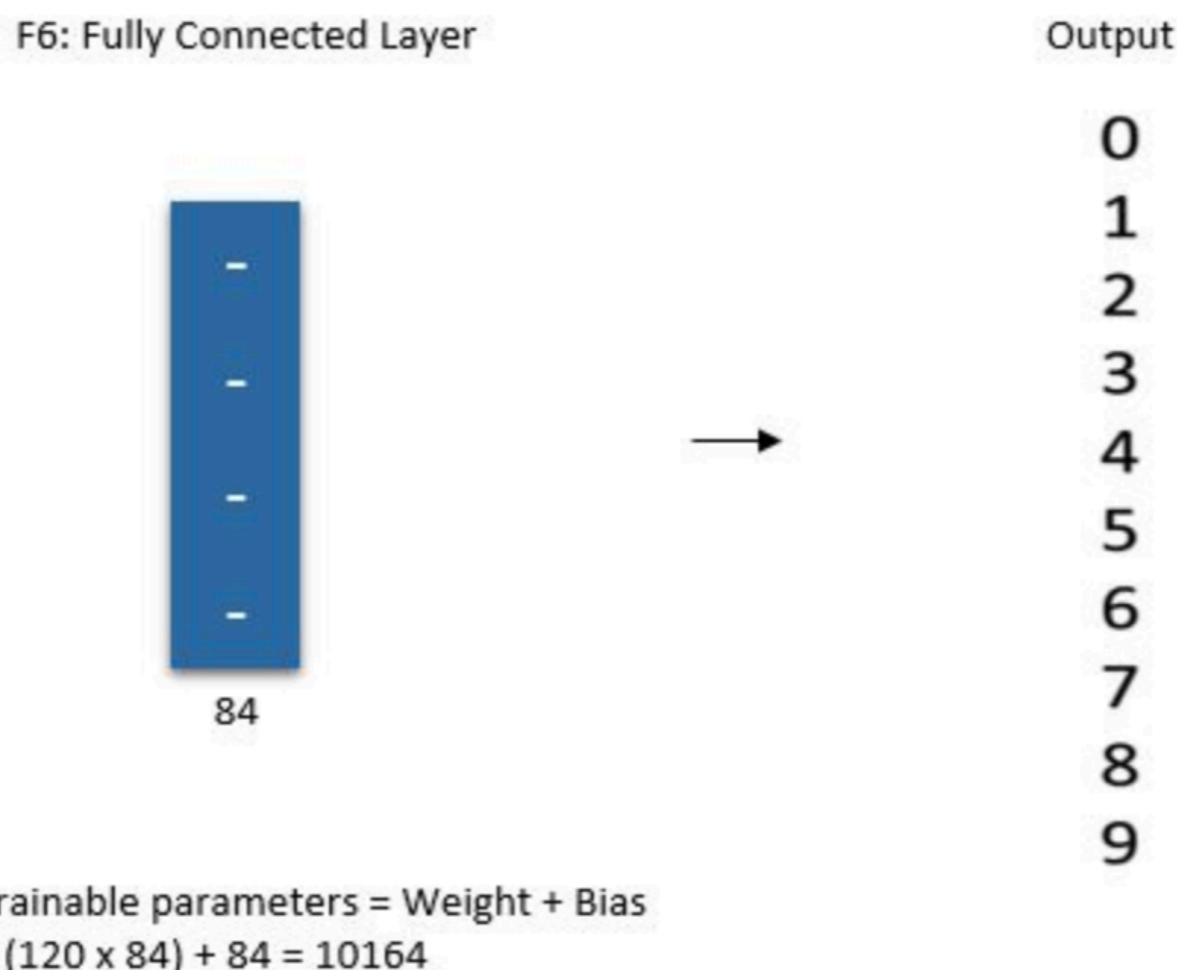
---



F6: Fully Connected Layer

# Some known architectures

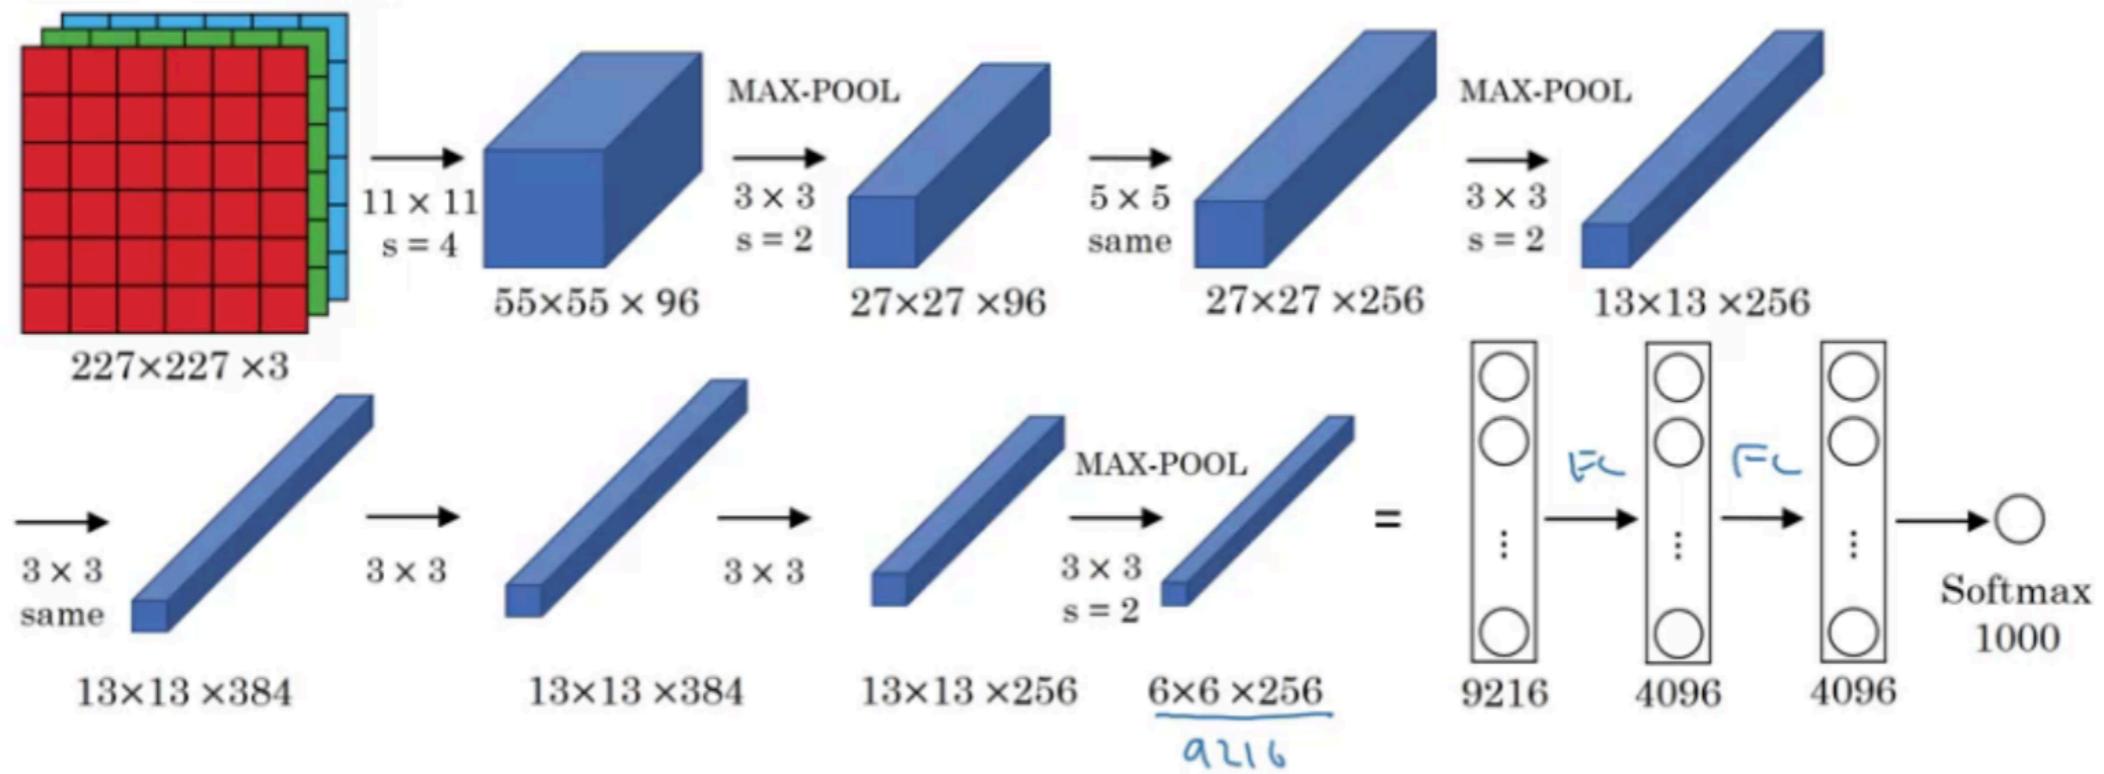
---



Fully Connected Output Layer

# Popular architectures

## AlexNet



This network is similar to LeNet-5 with just more convolution and pooling layers:

- **Parameters:** 60 million
- **Activation function:** ReLu